

Titanic Survival Prediction in python (Machine Learning project)

Required Libraries

```
In [7]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns
```

Step 1: Data frames Loading and importing

```
In [5]: df=pd.read_csv("C:/Users/DELL/Desktop/A/day 6(project)/train.csv")
```

```
In [6]: titanic_data = pd.read_csv("C:/Users/DELL/Desktop/A/day 6(project)/train.csv")
```

```
In [16...]: titanic_data.head()
```

Out[16...]

	PassengerId	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	1	0	3	male	22.0	1	0	7.2500	S
1	2	1	1	female	38.0	1	0	71.2833	C
2	3	1	3	female	26.0	0	0	7.9250	S
3	4	1	1	female	35.0	1	0	53.1000	S
4	5	0	3	male	35.0	0	0	8.0500	S

```
In [9]: titanic_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   PassengerId 891 non-null    int64  
 1   Survived     891 non-null    int64  
 2   Pclass       891 non-null    int64  
 3   Name         891 non-null    object  
 4   Sex          891 non-null    object  
 5   Age          714 non-null    float64 
 6   SibSp        891 non-null    int64  
 7   Parch        891 non-null    int64  
 8   Ticket       891 non-null    object  
 9   Fare          891 non-null    float64 
 10  Cabin         204 non-null    object  
 11  Embarked     889 non-null    object  
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

In [10...]: `titanic_data.describe()`

```
Out[10...]
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

shape of data sheet

```
In [11...]
```

```
titanic_data.shape
```

```
Out[11...]
```

```
(891, 12)
```

```
In [12...]
```

```
titanic_data.columns
```

```
Out[12...]
```

```
Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',  
       'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],  
      dtype='object')
```

Checking for Missing values

Step 2: Data Cleaning

```
In [13...]
```

```
titanic_data.isnull().sum()
```

```
Out[13...]: PassengerId      0  
          Survived        0  
          Pclass         0  
          Name          0  
          Sex           0  
          Age          177  
          SibSp         0  
          Parch         0  
          Ticket        0  
          Fare          0  
          Cabin        687  
          Embarked      2  
          dtype: int64
```


```
In [14...]: print(titanic_data.columns)  
  
Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',  
       'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],  
      dtype='object')
```

```
In [17...]: titanic_data = pd.read_csv("C:/Users/DELL/Desktop/A/day 6(project)/train.csv")  
  
# Now safely drop columns  
titanic_data.drop(['Ticket', 'Cabin', 'Name'], axis=1, inplace=True, errors='ignore')
```


```
In [21...]: import matplotlib.pyplot as plt  
import seaborn as sns  
  
f, ax = plt.subplots(1, 2, figsize=(12, 4))  
  
titanic_data['Survived'].value_counts().plot.pie(  
    explode=[0, 0.1], autopct='%1.1f%%', ax=ax[0], shadow=False)  
  
ax[0].set_title('Survivors (1) and the Dead (0)')  
ax[0].set_ylabel('')  
  
sns.countplot(x='Survived', data=titanic_data, ax=ax[1])
```

```
ax[1].set_ylabel('Quantity')
ax[1].set_title('Survivors (1) and the Dead (0)')
plt.show()
```

Survivors (1) and the Dead (0)




Survivors (1) and the Dead (0)



```
In [23...]: f, ax = plt.subplots(1, 2, figsize=(12, 4))
```

```
# Plot mean survival rate by gender
titanic_data[['Sex', 'Survived']].groupby(['Sex']).mean().plot.bar(ax=ax[0])
ax[0].set_title('Survivors by Sex')

# Countplot with hue for Survived
sns.countplot(x='Sex', hue='Survived', data=titanic_data, ax=ax[1])
ax[1].set_ylabel('Quantity')
ax[1].set_title('Survived (1) vs Dead (0): Men vs Women')
plt.tight_layout()
plt.show()
```



Step 3 : Feature Engineering (Optimizing Data for Model Training) next step

```
In [9]: train = pd.read_csv(r"C:\Users\DELL\Desktop\A\day 6(project)\train.csv")
test = pd.read_csv(r"C:\Users\DELL\Desktop\A\day 6(project)\test.csv")
```

```
In [ ]: # Drop unnecessary columns
train = train.drop(['Ticket', 'Name'], axis=1)
test = test.drop(['Ticket', 'Name'], axis=1)
```

```
In [8]: train = pd.read_csv("C:/Users/DELL/Desktop/A/day 6(project)/train.csv")
test = pd.read_csv("C:/Users/DELL/Desktop/A/day 6(project)/test.csv")
```

```
In [10...]: train = train.drop(['Ticket', 'Name'], axis=1)
test = test.drop(['Ticket', 'Name'], axis=1)
```

```
In [11...]: test
```

Out[11...]

	PassengerId	Pclass	Sex	Age	SibSp	Parch	Fare	Cabin	Embarked
0	892	3	male	34.5	0	0	7.8292	NaN	Q
1	893	3	female	47.0	1	0	7.0000	NaN	S
2	894	2	male	62.0	0	0	9.6875	NaN	Q
3	895	3	male	27.0	0	0	8.6625	NaN	S
4	896	3	female	22.0	1	1	12.2875	NaN	S
...
413	1305	3	male	NaN	0	0	8.0500	NaN	S
414	1306	1	female	39.0	0	0	108.9000	C105	C
415	1307	3	male	38.5	0	0	7.2500	NaN	S
416	1308	3	male	NaN	0	0	8.0500	NaN	S
417	1309	3	male	NaN	1	1	22.3583	NaN	C

418 rows × 9 columns

In [12...]

train

Out[12...]

	PassengerId	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Cabin	Embarked
0	1	0	3	male	22.0	1	0	7.2500	NaN	S
1	2	1	1	female	38.0	1	0	71.2833	C85	C
2	3	1	3	female	26.0	0	0	7.9250	NaN	S
3	4	1	1	female	35.0	1	0	53.1000	C123	S
4	5	0	3	male	35.0	0	0	8.0500	NaN	S
...
886	887	0	2	male	27.0	0	0	13.0000	NaN	S
887	888	1	1	female	19.0	0	0	30.0000	B42	S
888	889	0	3	female	NaN	1	2	23.4500	NaN	S
889	890	1	1	male	26.0	0	0	30.0000	C148	C
890	891	0	3	male	32.0	0	0	7.7500	NaN	Q

891 rows × 10 columns

In [13...]

train.isnull().sum()

```
Out[13...]: PassengerId      0  
          Survived        0  
          Pclass         0  
          Sex            0  
          Age           177  
          SibSp          0  
          Parch          0  
          Fare           0  
          Cabin          687  
          Embarked       2  
          dtype: int64
```

```
In [14...]: test.isnull().sum()
```

```
Out[14...]: PassengerId      0  
          Pclass         0  
          Sex            0  
          Age           86  
          SibSp          0  
          Parch          0  
          Fare           1  
          Cabin          327  
          Embarked       0  
          dtype: int64
```

```
In [16...]: train['Age'] = train['Age'].fillna(train['Age'].median())  
test['Age'] = test['Age'].fillna(test['Age'].median())  
  
train['Embarked'] = train['Embarked'].fillna(train['Embarked'].mode()[0])  
test['Fare'] = test['Fare'].fillna(test['Fare'].median())
```

```
In [20...]: train['Embarked'] = train['Embarked'].fillna(train['Embarked'].mode()[0])
```

```
In [21...]: test
```

Out[21...]

	PassengerId	Pclass	Sex	Age	SibSp	Parch	Fare	Cabin	Embarked
0	892	3	male	34.5	0	0	7.8292	NaN	Q
1	893	3	female	47.0	1	0	7.0000	NaN	S
2	894	2	male	62.0	0	0	9.6875	NaN	Q
3	895	3	male	27.0	0	0	8.6625	NaN	S
4	896	3	female	22.0	1	1	12.2875	NaN	S
...
413	1305	3	male	27.0	0	0	8.0500	NaN	S
414	1306	1	female	39.0	0	0	108.9000	C105	C
415	1307	3	male	38.5	0	0	7.2500	NaN	S
416	1308	3	male	27.0	0	0	8.0500	NaN	S
417	1309	3	male	27.0	1	1	22.3583	NaN	C

418 rows × 9 columns

In [22...]

train

Out[22...]

	PassengerId	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Cabin	Embarked
0	1	0	3	male	22.0	1	0	7.2500	NaN	S
1	2	1	1	female	38.0	1	0	71.2833	C85	C
2	3	1	3	female	26.0	0	0	7.9250	NaN	S
3	4	1	1	female	35.0	1	0	53.1000	C123	S
4	5	0	3	male	35.0	0	0	8.0500	NaN	S
...
886	887	0	2	male	27.0	0	0	13.0000	NaN	S
887	888	1	1	female	19.0	0	0	30.0000	B42	S
888	889	0	3	female	28.0	1	2	23.4500	NaN	S
889	890	1	1	male	26.0	0	0	30.0000	C148	C
890	891	0	3	male	32.0	0	0	7.7500	NaN	Q

891 rows × 10 columns

In [23...]

```
train = pd.get_dummies(train, columns=['Sex', 'Embarked'], drop_first=True)
test = pd.get_dummies(test, columns=['Sex', 'Embarked'], drop_first=True)
```

In [25...]

```
print("Train columns:", train.columns)
print("Test columns:", test.columns)
```

Train columns: Index(['PassengerId', 'Survived', 'Pclass', 'Age', 'SibSp', 'Parch', 'Fare',
 'Cabin', 'Sex_male', 'Embarked_Q', 'Embarked_S'],
 dtype='object')

Test columns: Index(['PassengerId', 'Pclass', 'Age', 'SibSp', 'Parch', 'Fare', 'Cabin',
 'Sex_male', 'Embarked_Q', 'Embarked_S'],
 dtype='object')

```
In [26...]: combined = pd.concat([train, test], sort=False)
```

```
In [28...]: combined = pd.concat([train, test], sort=False)
```

```
In [29...]: train = combined.loc[train.index]
test = combined.loc[test.index]
```

```
In [30...]: test
```

Out[30...]:

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare	Cabin	Sex_male	Embarked_Q	Embarked_S
0	1	0.0	3	22.0	1	0	7.2500	NaN	True	False	True
0	892	NaN	3	34.5	0	0	7.8292	NaN	True	True	False
1	2	1.0	1	38.0	1	0	71.2833	C85	False	False	False
1	893	NaN	3	47.0	1	0	7.0000	NaN	False	False	True
2	3	1.0	3	26.0	0	0	7.9250	NaN	False	False	True
...
415	1307	NaN	3	38.5	0	0	7.2500	NaN	True	False	True
416	417	1.0	2	34.0	1	1	32.5000	NaN	False	False	True
416	1308	NaN	3	27.0	0	0	8.0500	NaN	True	False	True
417	418	1.0	2	18.0	0	2	13.0000	NaN	False	False	True
417	1309	NaN	3	27.0	1	1	22.3583	NaN	True	False	False

836 rows × 11 columns

```
In [31...]: train
```

Out[31...]

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare	Cabin	Sex_male	Embarked_Q	Embarked_S
0	1	0.0	3	22.0	1	0	7.2500	NaN	True	False	True
0	892	NaN	3	34.5	0	0	7.8292	NaN	True	True	False
1	2	1.0	1	38.0	1	0	71.2833	C85	False	False	False
1	893	NaN	3	47.0	1	0	7.0000	NaN	False	False	True
2	3	1.0	3	26.0	0	0	7.9250	NaN	False	False	True
...
886	887	0.0	2	27.0	0	0	13.0000	NaN	True	False	True
887	888	1.0	1	19.0	0	0	30.0000	B42	False	False	True
888	889	0.0	3	28.0	1	2	23.4500	NaN	False	False	True
889	890	1.0	1	26.0	0	0	30.0000	C148	True	False	False
890	891	0.0	3	32.0	0	0	7.7500	NaN	True	True	False

1309 rows × 11 columns

In [32...]

```
X_train = train.drop(['PassengerId', 'Survived', 'Cabin'], axis=1)
y_train = train['Survived']

X_test = test.drop(['PassengerId', 'Cabin'], axis=1)
```

```
from sklearn.linear_model import LogisticRegression from sklearn.metrics import accuracy_score model =
LogisticRegression(max_iter=1000) model.fit(X_train, y_train) # Predict on train to check accuracy train_preds =
model.predict(X_train) print("Training Accuracy:", accuracy_score(y_train, train_preds))
```

In [34...]

```
print(train['Survived'].isnull().sum())
```

```
In [37...]: print(" train",train.head()) # Check what data you have in 'train'  
print("test",test.head()) # Check test data
```

```
train  PassengerId  Survived  Pclass  Age  SibSp  Parch  Fare Cabin  Sex_male  \  
0      1      0.0      3  22.0      1      0  7.2500  NaN    True  
0     892      NaN      3  34.5      0      0  7.8292  NaN    True  
1      2      1.0      1  38.0      1      0  71.2833  C85   False  
1     893      NaN      3  47.0      1      0  7.0000  NaN   False  
2      3      1.0      3  26.0      0      0  7.9250  NaN   False
```

```
Embarked_Q  Embarked_S  
0    False      True  
0    True      False  
1    False      False  
1    False      True  
2    False      True
```

```
test  PassengerId  Survived  Pclass  Age  SibSp  Parch  Fare Cabin  Sex_male  \  
0      1      0.0      3  22.0      1      0  7.2500  NaN    True  
0     892      NaN      3  34.5      0      0  7.8292  NaN    True  
1      2      1.0      1  38.0      1      0  71.2833  C85   False  
1     893      NaN      3  47.0      1      0  7.0000  NaN   False  
2      3      1.0      3  26.0      0      0  7.9250  NaN   False
```

```
Embarked_Q  Embarked_S  
0    False      True  
0    True      False  
1    False      False  
1    False      True  
2    False      True
```

```
In [38...]: train = pd.read_csv(r"C:\Users\DELL\Desktop\A\day 6(project)\train.csv")  
test = pd.read_csv(r"C:\Users\DELL\Desktop\A\day 6(project)\test.csv")
```

```
In [44...]: print("Train Survived nulls:", train['Survived'].isnull().sum()) # should be 0  
print("Test columns:", test.columns) # Should NOT have Survived column
```

```
Train Survived nulls: 0
Test columns: Index(['PassengerId', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp', 'Parch',
       'Ticket', 'Fare', 'Cabin', 'Embarked'],
      dtype='object')
```

```
In [45...]: combined = pd.concat([train.drop(columns=['Survived']), test], ignore_index=True)
# Do your feature engineering on combined here, then split later.
```

```
In [46...]: train_processed = combined.iloc[:len(train), :]
test_processed = combined.iloc[len(train):, :]
y_train = train['Survived'] # Keep target separate
```

```
In [47...]: X_train = train_processed
y_train = train['Survived']
X_test = test_processed
```

```
In [48...]: print(X_train.isnull().sum().sum()) # should be 0
print(X_test.isnull().sum().sum()) # should be 0
print(y_train.isnull().sum()) # should be 0
```

```
866
414
0
```

```
In [49...]: print("X_train missing values:")
print(X_train.isnull().sum()[X_train.isnull().sum() > 0])

print("\nX_test missing values:")
print(X_test.isnull().sum()[X_test.isnull().sum() > 0])
```

```
X_train missing values:
```

```
Age      177  
Cabin    687  
Embarked  2  
dtype: int64
```

```
X_test missing values:
```

```
Age      86  
Fare     1  
Cabin   327  
dtype: int64
```

```
In [56...]: X_train.drop(columns=['Cabin'], inplace=True, errors='ignore')  
X_test.drop(columns=['Cabin'], inplace=True, errors='ignore')
```

```
In [54...]: pd.options.mode.chained_assignment = None # 🤡 disables the warning
```

```
In [57...]: print("Cabin" in X_train.columns) # False = already gone  
False
```

```
In [59...]: # Drop columns not useful for model and not numerical  
X_train = X_train.select_dtypes(include=['number', 'bool'])  
X_test = X_test.select_dtypes(include=['number', 'bool'])
```

Model Training

```
In [61...]: from sklearn.linear_model import LogisticRegression  
from sklearn.metrics import accuracy_score  
  
model = LogisticRegression(max_iter=1000)  
model.fit(X_train, y_train)  
  
train_preds = model.predict(X_train)  
print("Train Accuracy:", accuracy_score(y_train, train_preds))
```

```
Train Accuracy: 0.7025813692480359
```

```
In [60...]: model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)
```

```
Out[60...]: LogisticRegression
```

```
In [63...]: from sklearn.model_selection import cross_val_score
scores = cross_val_score(model, X_train, y_train, cv=5)
print("Cross-validated accuracy:", scores.mean())
```

```
Cross-validated accuracy: 0.6891720544849664
```

```
In [64...]: test_preds = model.predict(X_test)
submission = pd.DataFrame({
    'PassengerId': test['PassengerId'],
    'Survived': test_preds
})
submission.to_csv('submission.csv', index=False)
```

```
In [71...]: import os

folder_path = r"D:\ML Projects\Titanic2"
os.makedirs(folder_path, exist_ok=True)
submission_path = os.path.join(folder_path, "submission.csv")
submission.to_csv(submission_path, index=False)

print(f"Submission file saved successfully at:\n{submission_path}")
```

```
Submission file saved successfully at:
D:\ML Projects\Titanic2\submission.csv
```